

STATS 762 Assignment 4

Francis Tang, ID 240887036

Due: 12 June 2019

Packages

```
library(MASS)
library(klaR)
library(nnet)
library(reshape2)
library(ggplot2)

## Registered S3 methods overwritten by 'ggplot2':
##   method      from
##   [.quosures   rlang
##   c.quosures   rlang
##   print.quosures rlang

library(glmnet)
```

```
## Loading required package: Matrix
## Loading required package: foreach
## Loaded glmnet 2.0-18

library(splines)
```

Question 1

(a)

First we read data and only subset the positions contain RDM, RCM and LS:

```
#read data
fifa=read.csv("~/Desktop/STATS 762/Fifa2019.csv")
fifa0 <- fifa[,-1]
fifa1 <- subset(fifa0, Position %in% c("RDM", "RCM", "LS"))
fifa1$Position <- factor(fifa1$Position)
```

Now we first try to fit a multinomial regression:

```
#fit a multinomial regression for the fifa data
fifa.mn <- multinom(Position ~ ., data = fifa1)
```

```
## # weights:  108 (70 variable)
## initial  value 929.425996
## iter   10 value 518.531804
## iter   20 value 504.714000
## iter   30 value 502.813422
## iter   40 value 493.590840
## iter   50 value 473.086244
```

```
## iter 60 value 448.736984
## iter 70 value 435.141344
## iter 80 value 413.418365
## iter 90 value 411.689334
## final value 411.687916
## converged
```

Make predictions based on the multinomial regression model, the results are shown in the matrix:

```
#prediction
mn.pred=predict(fifa.mn,fifa1)
#confusion matrix
table(mn.pred,fifa1$Position)
```

```
##
## mn.pred  LS RCM RDM
##      LS  203  4  2
##      RCM   4 320 123
##      RDM   0  67 123
```

Now let's try to fit LDA:

```
#Fit the LDA for the data
lda.fifa <- lda(Position ~ ., fifa1)
```

Make predictions based on the LDA model, the results are shown in the matrix:

```
#prediction
lda.pred=predict(lda.fifa,fifa1)
#confusion matrix
table(lda.pred$class,fifa1$Position)
```

```
##
##      LS RCM RDM
## LS  195  8  3
## RCM  12 310 116
## RDM   0  73 129
```

Now let's try to fit QDA:

```
#Fit the QDA for the train data
qda.fifa <- qda(Position ~ .,fifa1)
```

Make predictions based on the QDA model, the results are shown in the matrix:

```
#prediction
qda.pred=predict(qda.fifa,fifa1)
#confusion matrix
table(qda.pred$class,fifa1$Position)
```

```
##
##      LS RCM RDM
## LS  203  8  3
## RCM   4 336  63
## RDM   0  47 182
```

It is very obvious that QDA model achieved the best prediction accuracy. So in this case, we will pick QDA as our best model.

(b)

Here we need to find those rows which result in predicting RDM and RCM. This means that the membership probability of LS must be minimised and RDM and RCM need to be as close as possible. In this case, we pick 0.001 as the threshold for LS, [0.4,0.6] for RDM and RCM.

```
#print the class membership probability
qda.pred.prob.df <- as.data.frame.matrix(qda.pred$posterior)
#extract those rows which not resulting in LS
new1.df <- qda.pred.prob.df[(qda.pred.prob.df$LS < 0.001), ]
#one more step to extract those who result in both RDM and RCM (both probabilities between 0.4 to 0.6)
new2.df <- new1.df[new1.df$RCM > 0.4 & new1.df$RDM > 0.4, ]
new2.df
```

##		LS	RCM	RDM
##	1260	2.189597e-13	0.4275358	0.5724642
##	1306	7.801973e-16	0.5987345	0.4012655
##	1949	5.470421e-19	0.5218050	0.4781950
##	1960	4.135051e-13	0.5030562	0.4969438
##	2098	1.823022e-14	0.4633033	0.5366967
##	2458	1.434826e-16	0.4957068	0.5042932
##	3598	1.708051e-09	0.5496794	0.4503206
##	3683	1.536073e-18	0.5092774	0.4907226
##	3718	1.417468e-10	0.4497646	0.5502354
##	4200	2.356951e-17	0.5214958	0.4785042
##	4305	3.334352e-26	0.4386582	0.5613418
##	4539	1.872576e-27	0.4692426	0.5307574
##	4608	2.440561e-22	0.5997135	0.4002865
##	4866	4.648895e-19	0.5436244	0.4563756
##	4999	9.849763e-07	0.5158382	0.4841608
##	5702	1.188624e-16	0.4997141	0.5002859
##	5743	6.483179e-07	0.5769393	0.4230601
##	5796	2.001461e-18	0.4865715	0.5134285
##	5907	3.065174e-09	0.5142871	0.4857129
##	7238	3.005751e-09	0.5502656	0.4497344
##	7510	9.734893e-08	0.5449031	0.4550968
##	7985	6.256405e-09	0.5033189	0.4966811
##	8685	3.904366e-07	0.5245165	0.4754831
##	9418	2.342970e-07	0.5658330	0.4341667
##	9665	1.915169e-17	0.4174697	0.5825303
##	10163	4.861618e-07	0.5429784	0.4570212
##	10232	7.731262e-15	0.4842512	0.5157488
##	10419	1.435424e-22	0.4796220	0.5203780
##	10581	3.109864e-23	0.4835385	0.5164615
##	10788	5.123872e-08	0.4716201	0.5283798
##	11030	7.123522e-13	0.4193666	0.5806334
##	11374	1.819433e-15	0.5789556	0.4210444
##	12681	3.742474e-16	0.5597038	0.4402962
##	12786	2.062864e-07	0.5979880	0.4020118
##	12823	1.257851e-11	0.5777964	0.4222036
##	12888	5.243587e-16	0.4942304	0.5057696
##	13172	3.071227e-07	0.4294524	0.5705473
##	13382	1.598966e-15	0.5223960	0.4776040
##	14334	2.052584e-10	0.5259151	0.4740849
##	15094	1.647273e-11	0.4060602	0.5939398

```
## 15947 3.627611e-08 0.5010714 0.4989285
```

```
#match those performance scores in the original dataset and print out  
subset(fifa1[, -1], rownames(fifa1) %in% rownames(new2.df))
```

##	Crossing	Finishing	HeadingAccuracy	ShortPassing	Volley	Dribbling
## 1260	70	64	41	78	68	73
## 1306	62	56	63	80	63	72
## 1949	52	47	50	71	40	69
## 1960	78	68	61	80	62	73
## 2098	42	65	73	76	71	64
## 2458	60	56	68	76	50	75
## 3598	73	59	70	74	57	71
## 3683	63	49	47	68	71	68
## 3718	59	63	48	75	48	67
## 4200	59	53	58	74	56	71
## 4305	58	37	58	69	46	63
## 4539	50	39	54	72	38	67
## 4608	42	59	51	79	31	69
## 4866	53	48	55	73	57	64
## 4999	73	66	49	71	63	67
## 5702	44	55	62	68	39	65
## 5743	66	61	60	71	60	69
## 5796	46	49	58	64	35	58
## 5907	59	59	46	70	49	69
## 7238	67	55	64	67	51	62
## 7510	59	63	55	68	60	63
## 7985	52	59	65	75	63	64
## 8685	62	61	66	63	56	65
## 9418	66	58	58	68	64	66
## 9665	48	46	61	72	43	66
## 10163	57	54	48	67	45	70
## 10232	54	56	55	68	56	64
## 10419	41	41	51	65	41	58
## 10581	41	33	60	69	41	61
## 10788	61	52	52	67	57	65
## 11030	53	45	57	68	50	63
## 11374	49	41	68	63	46	56
## 12681	41	35	65	60	39	51
## 12786	53	54	41	64	47	59
## 12823	49	46	53	68	42	72
## 12888	50	48	61	63	46	50
## 13172	57	59	66	62	53	58
## 13382	61	29	45	54	35	62
## 14334	44	44	53	63	41	58
## 15094	48	43	45	67	46	59
## 15947	48	49	55	56	44	56
##	Curve	FKAccuracy	LongPassing	BallControl	Acceleration	SprintSpeed
## 1260	60	70	74	76	69	69
## 1306	69	67	76	76	67	66
## 1949	45	40	67	71	71	69
## 1960	83	85	79	78	55	45
## 2098	53	71	77	70	74	76
## 2458	42	45	70	77	67	72
## 3598	61	73	73	74	55	60

## 3683	68	68	70	72	83	74	
## 3718	58	63	68	70	69	75	
## 4200	66	56	70	74	61	60	
## 4305	53	50	67	67	66	63	
## 4539	42	41	69	71	70	72	
## 4608	48	38	77	69	47	47	
## 4866	49	48	66	68	64	65	
## 4999	68	65	67	70	78	77	
## 5702	53	39	67	67	47	41	
## 5743	64	58	68	72	74	68	
## 5796	47	42	60	63	63	66	
## 5907	55	60	65	71	74	68	
## 7238	59	62	65	64	76	72	
## 7510	59	66	66	68	66	63	
## 7985	66	65	68	68	52	54	
## 8685	55	56	60	65	76	73	
## 9418	68	68	61	67	65	64	
## 9665	48	43	61	70	67	60	
## 10163	64	54	64	67	68	68	
## 10232	58	53	63	67	64	60	
## 10419	38	41	59	64	54	41	
## 10581	37	40	66	69	61	68	
## 10788	62	57	59	72	70	60	
## 11030	48	45	63	67	71	65	
## 11374	48	29	58	60	48	44	
## 12681	48	49	57	57	67	66	
## 12786	46	43	63	62	70	71	
## 12823	47	35	62	69	79	66	
## 12888	56	45	62	56	54	58	
## 13172	62	60	58	62	57	54	
## 13382	54	36	50	58	77	76	
## 14334	48	38	58	62	64	62	
## 15094	49	41	61	61	59	62	
## 15947	49	44	55	57	60	64	
##	Agility	Reactions	Balance	ShotPower	Jumping	Stamina	Strength
## 1260	71	76	61	70	58	78	66
## 1306	70	73	66	66	53	80	69
## 1949	70	66	75	63	85	90	78
## 1960	58	70	61	78	68	66	65
## 2098	68	82	68	82	68	77	69
## 2458	68	73	67	73	68	76	76
## 3598	63	71	59	77	76	68	75
## 3683	91	74	92	76	82	86	36
## 3718	72	70	71	72	66	88	70
## 4200	64	73	69	73	82	72	67
## 4305	66	68	71	68	68	82	71
## 4539	69	66	62	65	63	77	66
## 4608	64	70	65	73	63	78	71
## 4866	71	67	68	65	71	79	72
## 4999	82	67	67	76	59	80	66
## 5702	58	67	67	67	68	81	78
## 5743	78	64	70	64	64	68	74
## 5796	63	64	62	53	53	94	76
## 5907	69	70	63	71	71	84	72

## 7238	71	66	64	62	73	82	66
## 7510	73	66	75	67	68	78	62
## 7985	61	64	58	73	61	68	65
## 8685	75	65	70	65	66	73	68
## 9418	73	61	72	65	79	67	74
## 9665	66	65	71	62	65	68	56
## 10163	80	61	77	61	74	68	64
## 10232	73	60	71	72	67	83	73
## 10419	59	71	63	64	58	76	66
## 10581	65	62	67	60	70	82	78
## 10788	72	61	74	59	49	70	64
## 11030	77	60	69	59	63	68	59
## 11374	34	60	45	51	67	69	81
## 12681	62	53	66	57	71	80	73
## 12786	71	62	73	59	63	66	67
## 12823	78	56	53	52	59	74	63
## 12888	50	61	49	65	52	71	80
## 13172	61	56	66	66	67	69	72
## 13382	73	45	67	52	68	86	60
## 14334	63	58	65	53	60	70	66
## 15094	62	55	60	52	58	58	67
## 15947	54	63	50	53	70	79	70
##	LongShots	Aggression	Interceptions	Positioning	Vision	Penalties	
## 1260	70	60		72	76	74	61
## 1306	64	71		76	55	76	51
## 1949	44	95		74	57	63	42
## 1960	78	60		62	63	80	84
## 2098	77	66		66	70	63	70
## 2458	68	80		73	66	71	49
## 3598	70	77		68	67	73	59
## 3683	78	63		74	68	68	57
## 3718	68	76		66	68	75	55
## 4200	66	73		72	61	66	57
## 4305	63	76		72	46	65	49
## 4539	59	74		70	47	70	47
## 4608	54	82		67	56	62	36
## 4866	55	80		70	49	60	48
## 4999	67	53		65	70	73	61
## 5702	63	81		64	55	65	63
## 5743	65	68		68	68	68	69
## 5796	54	77		70	55	60	44
## 5907	66	65		58	56	68	50
## 7238	54	67		70	63	65	68
## 7510	64	74		64	63	66	68
## 7985	70	69		66	65	69	55
## 8685	56	67		64	63	65	54
## 9418	59	59		63	66	72	65
## 9665	51	69		64	60	64	55
## 10163	54	54		60	59	65	48
## 10232	68	74		65	51	58	57
## 10419	58	64		64	54	57	40
## 10581	46	70		61	57	60	40
## 10788	45	60		58	66	67	47
## 11030	51	67		57	63	63	50

## 11374	39	66	61	49	55	39
## 12681	43	79	61	48	54	52
## 12786	53	62	61	52	66	49
## 12823	41	53	47	57	61	44
## 12888	62	67	54	41	57	58
## 13172	61	63	56	58	62	52
## 13382	43	56	56	42	49	44
## 14334	46	60	58	54	62	43
## 15094	39	57	52	61	61	48
## 15947	50	61	59	54	53	48
##	Composure	Marking	StandingTackle	SlidingTackle	GKDividing	GKHandling
## 1260	76	76	72	68	12	8
## 1306	70	71	73	68	12	13
## 1949	72	73	75	74	8	15
## 1960	75	70	65	58	9	7
## 2098	69	81	76	68	7	12
## 2458	70	68	70	66	15	8
## 3598	70	57	65	66	6	15
## 3683	67	70	66	62	10	8
## 3718	73	62	66	54	8	9
## 4200	72	72	73	67	9	9
## 4305	66	63	71	68	9	15
## 4539	68	66	70	67	9	6
## 4608	72	69	65	61	8	15
## 4866	69	64	69	61	8	9
## 4999	71	59	59	63	8	15
## 5702	64	62	70	66	7	8
## 5743	69	58	67	64	16	12
## 5796	67	69	70	68	12	12
## 5907	63	58	62	53	9	12
## 7238	64	70	64	63	14	10
## 7510	64	60	64	60	7	14
## 7985	66	65	59	56	15	14
## 8685	66	65	65	62	8	16
## 9418	72	53	61	54	10	9
## 9665	67	64	69	66	6	12
## 10163	62	60	62	58	8	15
## 10232	56	62	66	61	9	12
## 10419	66	58	62	58	11	10
## 10581	58	63	64	63	9	14
## 10788	62	56	55	52	8	11
## 11030	59	60	61	60	9	12
## 11374	52	63	64	59	7	8
## 12681	55	66	61	58	9	13
## 12786	51	60	61	60	11	8
## 12823	68	51	52	50	9	7
## 12888	53	59	65	62	12	6
## 13172	61	54	63	57	9	6
## 13382	55	59	58	59	12	8
## 14334	63	59	58	54	7	10
## 15094	59	58	56	55	8	11
## 15947	56	58	58	57	7	8
##	GKKicking	GKPositioning	GKReflexes			
## 1260	9	14	13			

## 1306	8	12	10
## 1949	9	16	11
## 1960	8	11	16
## 2098	13	7	8
## 2458	5	13	10
## 3598	9	15	7
## 3683	14	10	9
## 3718	14	15	15
## 4200	11	8	9
## 4305	14	10	8
## 4539	13	9	13
## 4608	14	9	6
## 4866	7	15	9
## 4999	16	10	14
## 5702	13	10	14
## 5743	12	11	16
## 5796	10	8	13
## 5907	15	14	7
## 7238	13	14	11
## 7510	8	7	9
## 7985	15	11	13
## 8685	12	8	14
## 9418	16	6	16
## 9665	9	5	12
## 10163	11	13	14
## 10232	8	9	12
## 10419	12	13	7
## 10581	14	12	12
## 10788	9	10	7
## 11030	10	7	12
## 11374	13	13	12
## 12681	13	13	9
## 12786	8	7	6
## 12823	8	11	8
## 12888	12	6	8
## 13172	10	6	9
## 13382	10	13	11
## 14334	10	11	5
## 15094	14	6	11
## 15947	14	13	8

(c)

We substitute the number given from the question then use QDA to predict the result:

```
qc.df = data.frame(Crossing = 57.487, Finishing = 57.71277, HeadingAccuracy = 58.64657,
  ShortPassing = 68.83688, Volleys = 54.40426, Dribbling = 65.74468,
  Curve = 57.09456, FKAAccuracy = 53.16312, LongPassing = 63.4539,
  BallControl = 68.76123, Acceleration = 67.00591, SprintSpeed = 66.63475,
  Agility = 68.67376, Reactions = 66.62648, Balance = 67.78369,
  ShotPower = 67.30378, Jumping = 67.24232, Stamina = 73.51773,
  Strength = 69.20331, LongShots = 61.43735, Aggression = 65.65839,
  Interceptions = 55.4669, Positioning = 62.02719, Vision = 63.98818,
```



```

Penalties = 57.40189, Composure = 65.89835, Marking = 54.90898,
StandingTackle = 55.4669, SlidingTackle = 51.90544, GKDividing = 10.69267,
GKHandling = 10.63357, GKKicking = 10.83333, GKPositioning = 10.65248,
GKReflexes = 10.69031)
qda.pred1=predict(qda.fifa,qc.df)
qda.pred1

## $class
## [1] RCM
## Levels: LS RCM RDM
##
## $posterior
##           LS           RCM           RDM
## 1 1.964279e-06 0.7651683 0.2348298

```

(d)

Here we fit a classification tree:

```

library(rpart); library(rpart.plot); library(rattle); library(gbm)

## Rattle: A free graphical interface for data science with R.
## Version 5.2.0 Copyright (c) 2006-2018 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.

## Loaded gbm 2.1.5
#Fit a classification tree
set.seed(1e5)
fifa.cart0 <- rpart(Position~., data=fifa1,method='class',cp=0.001)
fifa.cart0$cptable

```

```

##           CP nsplit rel error    xerror    xstd
## 1  0.364835165      0 1.0000000 1.0000000 0.03187113
## 2  0.053846154      1 0.6351648 0.6417582 0.03039133
## 3  0.017582418      3 0.5274725 0.6109890 0.03002617
## 4  0.015384615      5 0.4923077 0.6153846 0.03008089
## 5  0.013186813      6 0.4769231 0.6131868 0.03005364
## 6  0.008791209      7 0.4637363 0.6021978 0.02991415
## 7  0.007692308      9 0.4461538 0.6043956 0.02994248
## 8  0.007326007     11 0.4307692 0.6087912 0.02999849
## 9  0.006593407     14 0.4087912 0.6065934 0.02997060
## 10 0.004395604     17 0.3846154 0.6021978 0.02991415
## 11 0.003296703     22 0.3626374 0.6087912 0.02999849
## 12 0.003076923     24 0.3560440 0.6153846 0.03008089
## 13 0.002197802     30 0.3318681 0.6175824 0.03010792
## 14 0.001000000     33 0.3252747 0.6373626 0.03034169

```

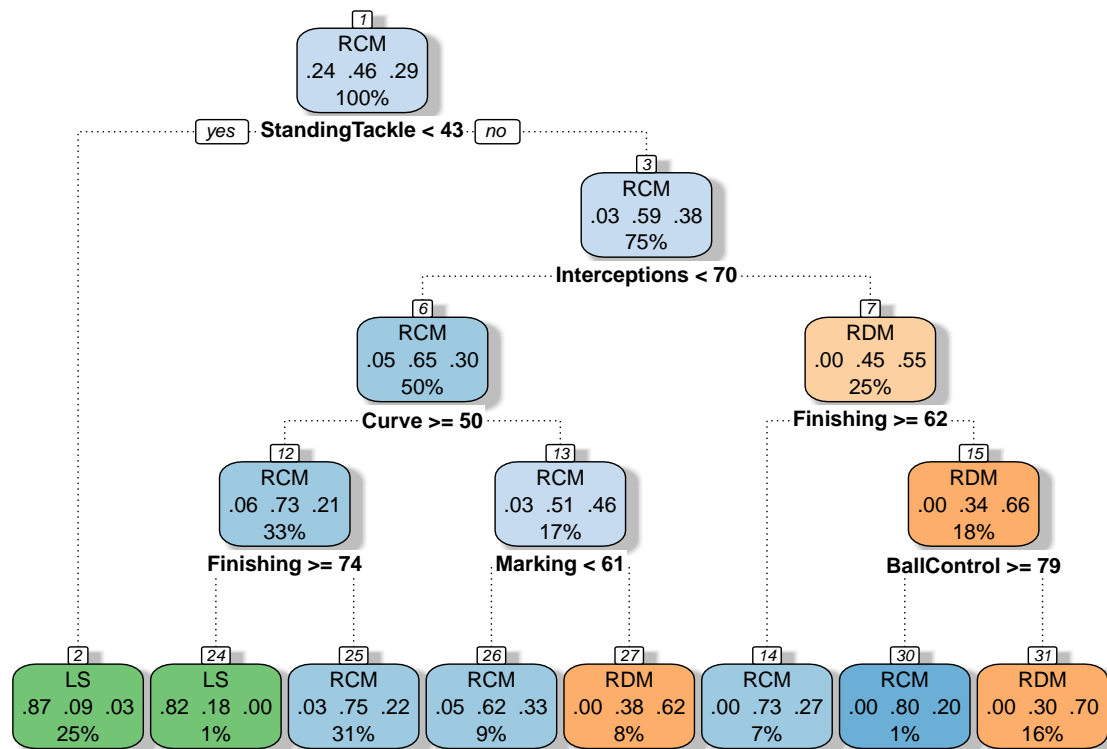
Now we prune the tree with a particular cp. We noticed that 0.008791209 has the smallest cross validation error 0.6021978, so we use this one to prune the tree. We will also compare this with the full model.

```

#Prune the tree with a particular complexity paramter (cp)

fifa.prune0 <-prune(fifa.cart0,cp=fifa.cart0$cptable[6,1])
fancyRpartPlot(fifa.prune0, uniform=TRUE,main=" ")

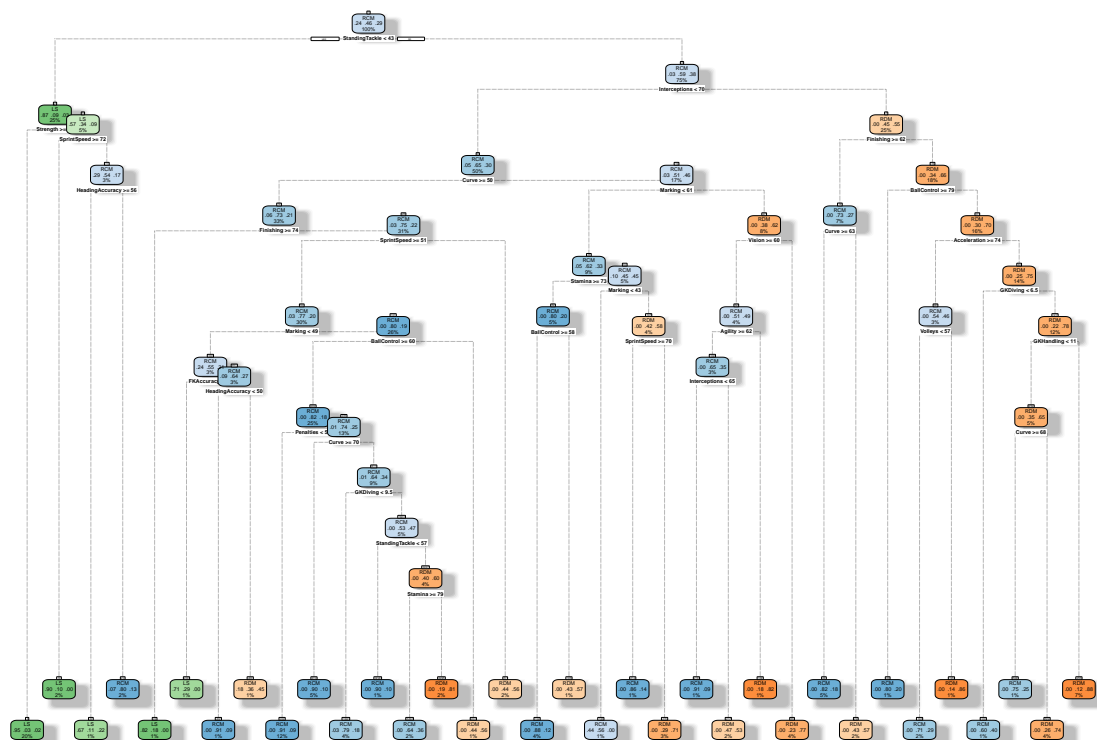
```



Rattle 2019-Jun-12 23:59:05 francistang

```
fifa.prune0.full <-prune(fifa.cart0,cp=fifa.cart0$cptable[14,1])
fancyRpartPlot(fifa.prune0.full, uniform=TRUE,main=" ")
```

```
## Warning: labs do not fit even at cex 0.15, there may be some overplotting
```



Rattle 2019-Jun-12 23:59:05 francistang

(e)

Now we use pruned tree to make predictions also comparing the full model. Although the full model has better predictions, it may suffer from the problem of overfitting.

```
fifa.pred0 <- predict(fifa.prune0,newdata=fifa1[,-1],type='class')
fifa.pred0.full <- predict(fifa.prune0.full,newdata=fifa1[,-1],type='class')
table(fifa.pred0,fifa1$Position)
```

```
##
## fifa.pred0  LS RCM RDM
##           LS 195 22  7
##           RCM 12 301 102
##           RDM  0  68 139
```

```
table(fifa.pred0.full,fifa1$Position)
```

```
##
## fifa.pred0.full  LS RCM RDM
##                LS 199 12  5
##                RCM  6 311 55
##                RDM  2  68 188
```

Comparing to the QDA model, the classification tree does not achieve a better prediction accuracy. So the best model will remain as the QDA model in (a).

```
#confusion matrix
table(qda.pred$class,fifa1$Position)
```

```
##
```

```
##          LS RCM RDM
##   LS   203   8   3
##   RCM   4 336  63
##   RDM   0  47 182
```

Question 2:

(a)

In this question, instead of using Position in Question 1, we use Overall for processing.

```
#keep Overall but get rid of Position
fifa3 <- subset(fifa, Position %in% c("RDM", "RCM", "LS"))
fifa3$Position <- factor(fifa3$Position)
fifa2 <- fifa3[,-2]
```

First, we fit a regression tree with $cp=0.001$

```
#Setup random numbers
set.seed(1e5)

#Fit a regression tree with cp=0.001
fifa2.cart <- rpart(Overall~. , data=fifa2,method='anova',cp=0.001)
fifa2.cart$cptable
```

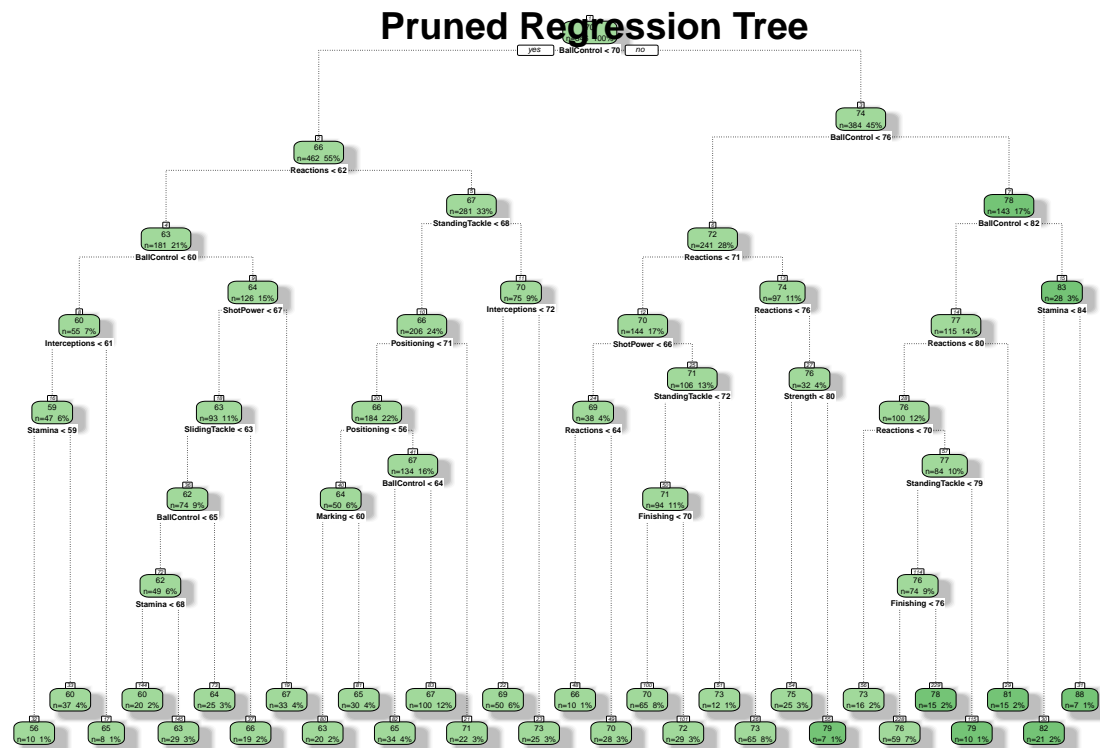
```
##          CP nsplit  rel error    xerror    xstd
## 1  0.492527507      0 1.00000000 1.0009109 0.05147726
## 2  0.106461467      1 0.50747249 0.5334657 0.03031492
## 3  0.078900624      2 0.40101103 0.4292515 0.02326115
## 4  0.033098990      3 0.32211040 0.3581378 0.01978659
## 5  0.025668595      4 0.28901141 0.3174626 0.01859831
## 6  0.025577783      5 0.26334282 0.2955118 0.01666606
## 7  0.022393129      6 0.23776503 0.2835535 0.01655685
## 8  0.016256616      7 0.21537190 0.2717535 0.01612412
## 9  0.010207695      8 0.19911529 0.2585760 0.01562668
## 10 0.009165922      9 0.18890759 0.2453542 0.01490193
## 11 0.008946992     10 0.17974167 0.2419035 0.01427666
## 12 0.006272348     11 0.17079468 0.2357554 0.01387268
## 13 0.005889702     12 0.16452233 0.2294103 0.01359778
## 14 0.005864620     13 0.15863263 0.2265236 0.01361185
## 15 0.005738690     14 0.15276801 0.2265236 0.01361185
## 16 0.005533938     15 0.14702932 0.2231265 0.01314817
## 17 0.005172275     16 0.14149538 0.2241645 0.01337391
## 18 0.004455127     17 0.13632311 0.2219853 0.01333352
## 19 0.003858016     18 0.13186798 0.2199174 0.01297966
## 20 0.003848973     19 0.12800996 0.2160672 0.01292953
## 21 0.003436386     20 0.12416099 0.2131558 0.01290195
## 22 0.003333182     21 0.12072461 0.2095240 0.01280840
## 23 0.003284612     22 0.11739142 0.2076158 0.01238015
## 24 0.002440085     23 0.11410681 0.2034690 0.01217081
## 25 0.002419914     25 0.10922664 0.2002660 0.01229456
## 26 0.002405192     26 0.10680673 0.1996279 0.01228839
## 27 0.002318690     27 0.10440154 0.1997131 0.01229406
## 28 0.001965770     29 0.09976416 0.1980072 0.01214966
```

```
## 29 0.001944549    30 0.09779839 0.1910233 0.01183062
## 30 0.001871194    31 0.09585384 0.1911958 0.01183486
## 31 0.001839383    32 0.09398264 0.1912890 0.01179788
## 32 0.001832131    33 0.09214326 0.1912890 0.01179788
## 33 0.001529021    34 0.09031113 0.1886008 0.01123284
## 34 0.001518805    35 0.08878211 0.1890735 0.01125954
## 35 0.001501875    36 0.08726330 0.1881291 0.01125729
## 36 0.001310566    37 0.08576143 0.1875399 0.01123321
## 37 0.001289367    38 0.08445086 0.1876222 0.01120510
## 38 0.001255896    39 0.08316150 0.1871208 0.01117898
## 39 0.001249564    40 0.08190560 0.1876353 0.01119065
## 40 0.001239840    41 0.08065604 0.1876353 0.01119065
## 41 0.001233527    42 0.07941620 0.1879520 0.01118441
## 42 0.001213148    43 0.07818267 0.1873484 0.01112875
## 43 0.001115453    45 0.07575637 0.1891926 0.01119668
## 44 0.001000000    46 0.07464092 0.1889483 0.01115605
```

When $\alpha = 0.001255896$, the cv-error is minimised: 0.1871208. $\alpha = 0.001965770$ is the largest value in which the corresponding cv-error: 0.1980072 is within the one standard deviation around the minimum error: $0.1871208 + 0.01117898 = 0.19829978$.

Now we prune the trees and plot it out together:

```
#Prune trees
fifa2.opt=prune(fifa2.cart,cp=fifa2.cart$cptable[28,1])
fancyRpartPlot(fifa2.opt, uniform=TRUE,main="Pruned Regression Tree")
```



Rattle 2019-Jun-12 23:59:06 francistang

```
fifa2.opt
```

```
## n= 846
```

```

##
## node), split, n, deviance, yval
##      * denotes terminal node
##
## 1) root 846 31079.27000 69.51655
##    2) BallControl< 69.5 462 7806.63400 65.63853
##      4) Reactions< 61.5 181 2490.25400 62.76796
##        8) BallControl< 59.5 55 722.80000 59.80000
##          16) Interceptions< 60.5 47 353.23400 58.87234
##            32) Stamina< 58.5 10 53.60000 55.80000 *
##            33) Stamina>=58.5 37 179.72970 59.70270 *
##          17) Interceptions>=60.5 8 91.50000 65.25000 *
##        9) BallControl>=59.5 126 1071.49200 64.06349
##          18) ShotPower< 66.5 93 611.69890 63.11828
##            36) SlidingTackle< 62.5 74 432.50000 62.50000
##              72) BallControl< 64.5 49 255.83670 61.59184
##                144) Stamina< 67.5 20 135.80000 60.10000 *
##                145) Stamina>=67.5 29 44.82759 62.62069 *
##              73) BallControl>=64.5 25 57.04000 64.28000 *
##            37) SlidingTackle>=62.5 19 40.73684 65.52632 *
##          19) ShotPower>=66.5 33 142.54550 66.72727 *
##        5) Reactions>=61.5 281 2864.20600 67.48754
##          10) StandingTackle< 67.5 206 1623.32500 66.47087
##            20) Positioning< 70.5 184 1052.08200 65.92935
##              40) Positioning< 55.5 50 284.50000 64.30000
##                80) Marking< 59.5 20 102.95000 62.55000 *
##                81) Marking>=59.5 30 79.46667 65.46667 *
##              41) Positioning>=55.5 134 585.31340 66.53731
##                82) BallControl< 63.5 34 174.97060 65.02941 *
##                83) BallControl>=63.5 100 306.75000 67.05000 *
##            21) Positioning>=70.5 22 66.00000 71.00000 *
##          11) StandingTackle>=67.5 75 443.12000 70.28000
##            22) Interceptions< 71.5 50 164.02000 69.14000 *
##            23) Interceptions>=71.5 25 84.16000 72.56000 *
##        3) BallControl>=69.5 384 7965.24000 74.18229
##          6) BallControl< 75.5 241 2273.50200 71.92116
##            12) Reactions< 70.5 144 969.30560 70.43056
##              24) ShotPower< 65.5 38 259.07890 68.60526
##                48) Reactions< 63.5 10 45.60000 65.80000 *
##                49) Reactions>=63.5 28 106.67860 69.60714 *
##              25) ShotPower>=65.5 106 538.23580 71.08491
##                50) StandingTackle< 71.5 94 417.15960 70.79787
##                  100) Finishing< 69.5 65 229.13850 70.16923 *
##                  101) Finishing>=69.5 29 104.75860 72.20690 *
##                51) StandingTackle>=71.5 12 52.66667 73.33333 *
##            13) Reactions>=70.5 97 509.25770 74.13402
##              26) Reactions< 75.5 65 207.53850 73.23077 *
##              27) Reactions>=75.5 32 140.96880 75.96875
##                54) Strength< 79.5 25 43.36000 75.16000 *
##                55) Strength>=79.5 7 22.85714 78.85714 *
##          7) BallControl>=75.5 143 2382.99300 77.99301
##            14) BallControl< 81.5 115 993.44350 76.66957
##              28) Reactions< 79.5 100 639.64000 76.06000
##              56) Reactions< 69.5 16 78.00000 73.00000 *


```

```
##          57) Reactions>=69.5 84   383.28570 76.64286
##          114) StandingTackle< 78.5 74   290.21620 76.32432
##          228) Finishing< 75.5 59   169.55930 75.79661 *
##          229) Finishing>=75.5 15    39.60000 78.40000 *
##          115) StandingTackle>=78.5 10    30.00000 79.00000 *
##          29) Reactions>=79.5 15    68.93333 80.73333 *
##          15) BallControl>=81.5 28   360.85710 83.42857
##          30) Stamina< 83.5 21   138.95240 81.95238 *
##          31) Stamina>=83.5 7    38.85714 87.85714 *
```

(b)

This question requires us to fit a gradient boosting regression tree:

```
#Fit a boosting reg tree
fifa.gbm <- gbm(Overall~., data = fifa2, distribution='gaussian',
               shrinkage = 0.04, n.trees = 4000, cv.folds = 10)
fifa.gbm.perf = gbm.perf(fifa.gbm, method = "cv")
```



```
fifa.gbm.perf
```

```
## [1] 3058
```

From the result above we can conclude that the optimal number of trees are 3058.

(c)

```
#Predict values and find the MSE using optimal regression tree
fifa2.opt.pred <- predict(fifa2.opt,newdata=fifa2[,-1],type='vector')
opt.res=fifa2.opt.pred-fifa2$Overall;
mean(opt.res^2)
```

```
## [1] 3.665008
#Predict values and find the MSE using optimal gradient boosting regression tree
fifa.gbm.pred <- predict(fifa.gbm,newdata = fifa2[,-1],n.trees = fifa.gbm.perf,type = "response")
#mse
fifa.res=fifa.gbm.pred-fifa2$Overall;
mean(fifa.res^2)

## [1] 1.097266
#Predict values and find the MSE using optimal linear regression with lasso
set.seed(1e5)
cv.lasso=cv.glmnet(as.matrix(fifa2[,-1]),as.matrix(fifa2[,1]),alpha=1,standardize=TRUE)
fifa.lasso.pred <- predict(cv.lasso, as.matrix(fifa2[,-1]),
type='response',lambda=cv.lasso$lambda.1se)
lasso.res = fifa.lasso.pred - fifa2$Overall
mse3 = mean(lasso.res^2)
mse3
```

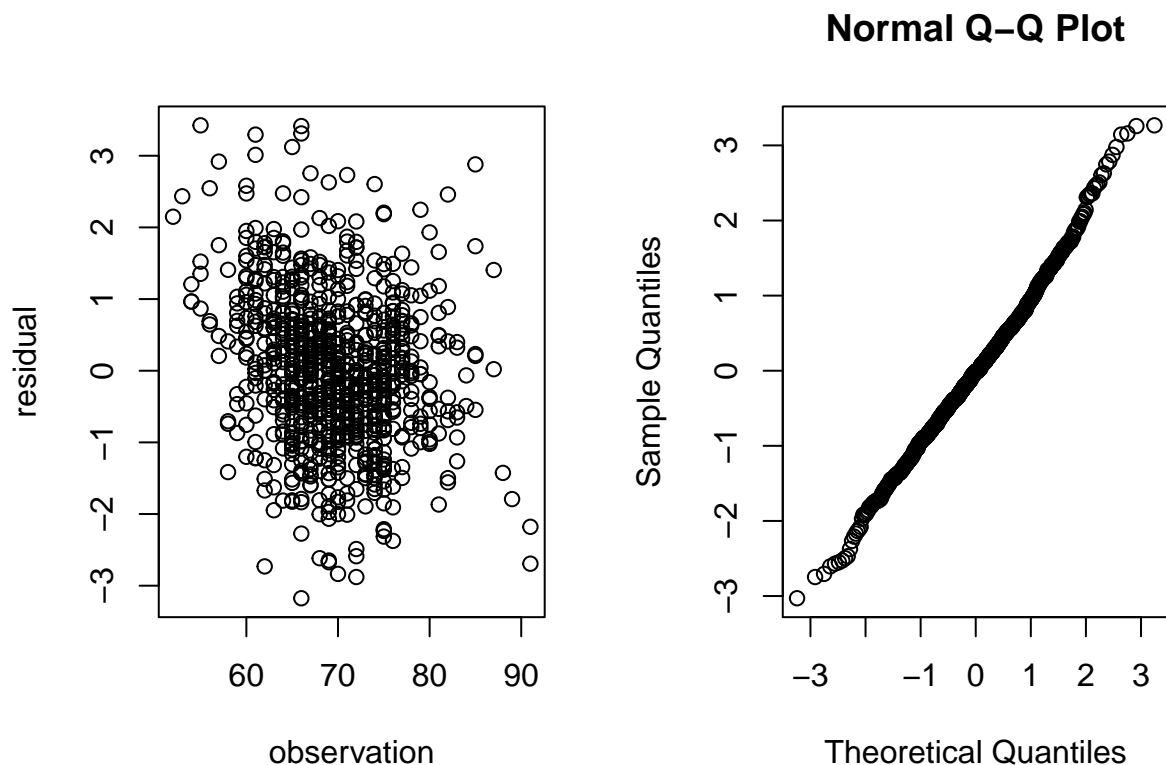
```
## [1] 3.909375
```

From the comparison of MSE among those three model above, we found out that the optimal gradient boosting regression tree model achieved the smallest MSE, so we are able to conclude that the optimal gradient boosting regression tree model is the best model for this case.

(d)

In this question, we are required to plot residual against overall score for the optimal gradient boosting regression tree:

```
par(mfrow=c(1,2)); plot(fifa2$Overall,fifa.res,xlab='observation',ylab='residual');
qqnorm(fifa.res/sd(fifa.res))
```

The residual VS observation plot shows no obvious pattern and most of the points concentrate between $[-2, 2]$ which is actually a good result for modelling. Normal Q-Q plots also proves the same conclusion as almost all points stand on the line together.


(e)

This question requires us to compare the relative variable importance of both my trees in (a) and (b).

```
#relative variable importance of optimal regression tree
fifa2.cart$variable.importance/sum(fifa2.cart$variable.importance)
```

##	BallControl	Reactions	Dribbling	ShortPassing
##	0.2305382257	0.1369057188	0.1318139069	0.1261103201
##	Vision	LongPassing	Positioning	StandingTackle
##	0.1134453557	0.0944542590	0.0201117160	0.0196792053
##	Interceptions	LongShots	Composure	Marking
##	0.0169788065	0.0151097339	0.0146184851	0.0122423988
##	SlidingTackle	Finishing	Aggression	ShotPower
##	0.0117693760	0.0090674022	0.0085215933	0.0081859016
##	Stamina	HeadingAccuracy	Crossing	Volleys
##	0.0055878179	0.0047447899	0.0040151015	0.0035855561
##	Strength	Penalties	Balance	Agility
##	0.0033779482	0.0021655553	0.0016112820	0.0011732980
##	Acceleration	FKAccuracy	Curve	GKHandling
##	0.0010376560	0.0009955272	0.0008877492	0.0004246820
##	SprintSpeed	Jumping	GKDividing	
##	0.0003646797	0.0002660874	0.0002098648	

```
#relative variable importance of optimal gradient boosting regression tree
fifa.gbm.summary <- summary.gbm(fifa.gbm)
```



In order to compare them, we combine them together to compare:

```
ort.rvi <- as.matrix(fifa2.cart$variable.importance/sum(fifa2.cart$variable.importance) * 100)
gbm.rvi <- fifa.gbm.summary$rel.inf
compare.rvi = cbind(ort.rvi, gbm.rvi[match(rownames(ort.rvi), rownames(fifa.gbm.summary))])
colnames(compare.rvi) = c("regression tree", "gradient boosting tree")
compare.rvi
```

##	regression tree	gradient boosting tree
## BallControl	23.05382257	33.03098084
## Reactions	13.69057188	28.47428063
## Dribbling	13.18139069	1.20486159
## ShortPassing	12.61103201	3.75838490
## Vision	11.34453557	0.33137879
## LongPassing	9.44542590	0.57950165
## Positioning	2.01117160	3.34531457
## StandingTackle	1.96792053	2.30487147
## Interceptions	1.69788065	1.17245649
## LongShots	1.51097339	0.63527299
## Composure	1.46184851	4.98385002
## Marking	1.22423988	0.82443586
## SlidingTackle	1.17693760	1.45527133
## Finishing	0.90674022	2.90463613
## Aggression	0.85215933	1.10960971
## ShotPower	0.81859016	4.74214808
## Stamina	0.55878179	1.10378380
## HeadingAccuracy	0.47447899	2.05773145
## Crossing	0.40151015	0.34733525
## Volleys	0.35855561	0.55332566
## Strength	0.33779482	0.37048659
## Penalties	0.21655553	0.39807456
## Balance	0.16112820	0.51466982

## Agility	0.11732980	0.42650598
## Acceleration	0.10376560	0.57380160
## FKAccuracy	0.09955272	0.45623188
## Curve	0.08877492	0.22790807
## GKHandling	0.04246820	0.13829641
## SprintSpeed	0.03646797	0.74744805
## Jumping	0.02660874	0.65216298
## GKDiving	0.02098648	0.09076354

To find out which variables have similar importance score, we allow them to have 10% difference in this case:

```
compare.rvi[0.9 < compare.rvi[,1]/compare.rvi[,2] & compare.rvi[,1]/compare.rvi[,2] < 1.11, ]
```

##	regression tree	gradient boosting tree
##	0.3377948	0.3704866

In this case we have “Strength” has similar results between tree models in (a) and (b). Let’s loose the boundary of similarity:

```
compare.rvi[0.75 < compare.rvi[,1]/compare.rvi[,2] & compare.rvi[,1]/compare.rvi[,2] < 1.33, ]
```

##	regression tree	gradient boosting tree
## StandingTackle	1.9679205	2.3048715
## SlidingTackle	1.1769376	1.4552713
## Aggression	0.8521593	1.1096097
## Crossing	0.4015102	0.3473353
## Strength	0.3377948	0.3704866

Now we have four more! “StandingTackle”, “SlidingTackle”, “Aggression”, “Crossing” and “Strength”. So we can conclude that these five are roughly equally important.

Now it comes with the question: why all variables are NOT equally important? To answer this question, we need to go back to what actually determines relative variable importances - it depends on how many times this variable has been chosen for splitting the tree, which means a variable achieves higher importance score only because it has been chosen for more time than others. Rather than just calculate the times when the variable got chosen for splitting for optimal regression trees, optimal gradient boosting regression tree sums its importance of each trees separately. The sum will be divided by the total number of trees, which makes it different from normal regression trees, so the difference comes out.

Question 3

(a)

First, we read the data and give indexes to the rows:

```
accident <- read.csv("~/Desktop/STATS 762/airliner_accidents-1.csv")
accident$index <- 1:72
```

The goal is to model the fatal accidents with year using both natural splines and B-splines to fit. We fit the temperature data using natural splines and B-splines with the degree of 3. Various number of inner knots are considered, 0-10.

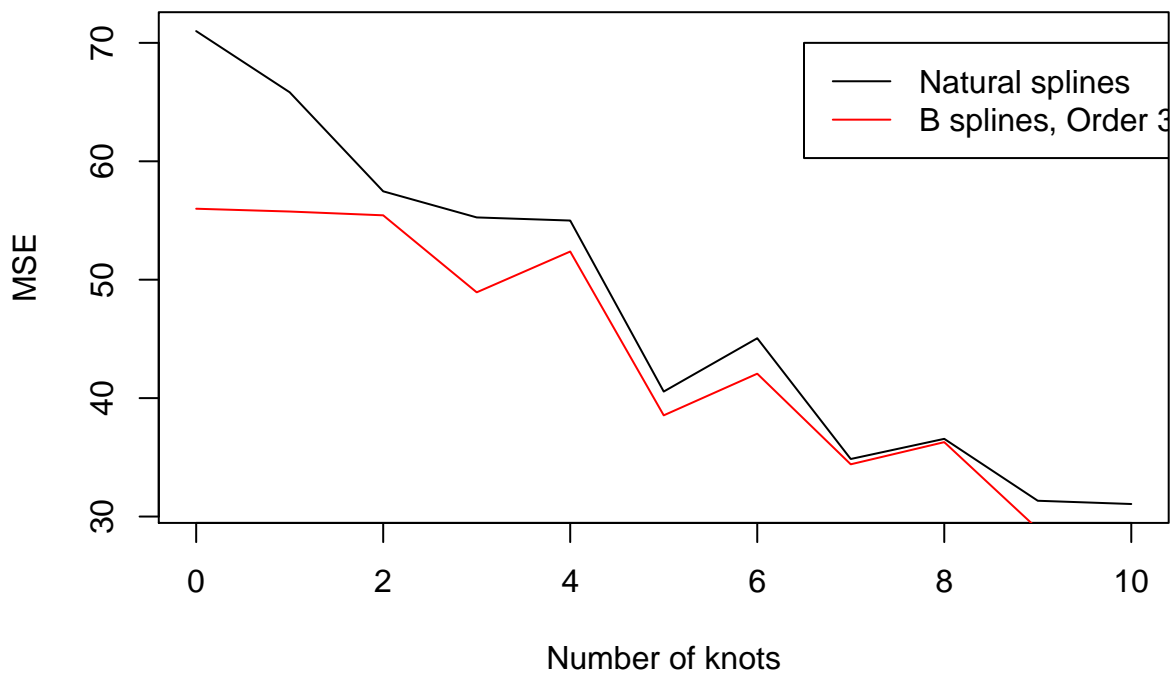
```
set.seed(1e5)
acci.mse.ns=acci.mse.bs=c(0:10)
n.knots=c(0:10)
for(j in 1:length(acci.mse.ns)){
```

```

#natural splines
a.ns=ns(accident$index,df=n.knots[j]+1,intercept=FALSE)
#B splines
a.bs=bs(accident$index,df=n.knots[j]+3,intercept=FALSE)
#predict accidents
pre.acci.ns=predict(lm(accident$Fatal~a.ns), interval='confidence');
pre.acci.bs=predict(lm(accident$Fatal~a.bs), interval='confidence');
#MSE
acci.mse.ns[j]=mean((accident$Fatal-pre.acci.ns[,1])^2)
acci.mse.bs[j]=mean((accident$Fatal-pre.acci.bs[,1])^2)
}

plot(n.knots,acci.mse.ns,'l',ylab='MSE',xlab='Number of knots'); lines(n.knots,acci.mse.bs,col='red');
legend(6.5,70,legend=c("Natural splines", "B splines, Order 3"),col=c(1,2),lty=1);

```



LOOCV error -vs- number of knots:


```

#LOOCV error -vs- number of knots
acci.cv.ns=acci.cv.bs=rep(0,length(n.knots))
for(j in 1:length(n.knots)){ for(l in 1:length(accident$index)){
  #predict accidents
  pre.a.ns=predict(lm(Fatal~ns(index,df=n.knots[j]+1,
    intercept=FALSE,Boundary.knots=c(1,72)),
    data=accident[-l,],newdata=accident[l,])
  pre.a.bs=predict(lm(Fatal~bs(index,df=n.knots[j]+3,intercept=FALSE,
    Boundary.knots=c(1,72)),data=accident[-l,],
    newdata=accident[l,])

  #cumulative sum of error
  acci.cv.ns[j]=acci.cv.ns[j]+(accident$Fatal[l]-pre.a.ns)^2
  acci.cv.bs[j]=acci.cv.bs[j]+(accident$Fatal[l]-pre.a.bs)^2
}}
acci.cv.ns=acci.cv.ns/length(accident$index)
acci.cv.bs=acci.cv.bs/length(accident$index)

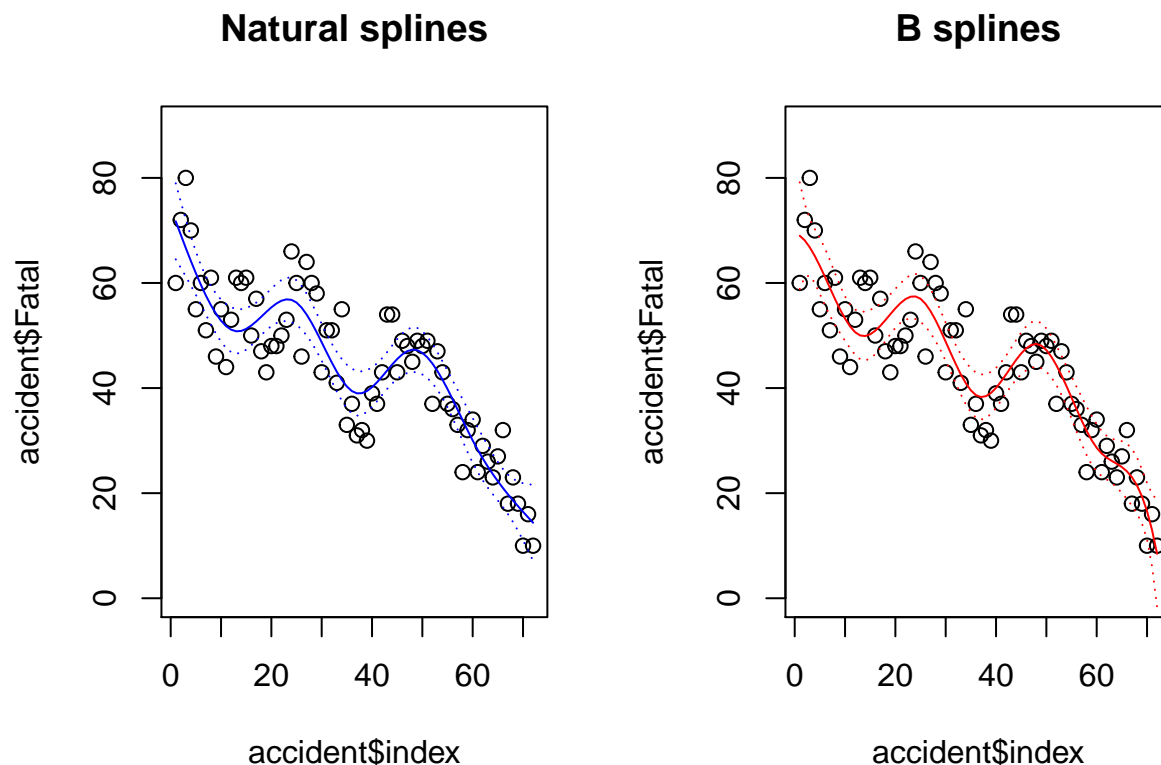
```

```
plot(n.knots,log(acci.cv.ns),ylab='log(cv error)',xlab='Number of knots','o',ylim=c(3.5,4.5)); lines(n.knots,log(acci.cv.bs),ylab='log(cv error)',xlab='Number of knots','o',ylim=c(3.5,4.5)); legend(6,4.5,legend=c("Natural splines","B splines, Order 3"),col=c(1,2), lty=1);
```



It is not hard to figure out that when the natural splines has the smallest error when number of knots equals 7 and B splines has the smallest error when number of knots equals to 10.

```
acci0.ns=predict(lm(Fatal~ns(index,df=5+1,intercept=FALSE),data=accident),
                 newdata=accident,interval = 'confidence')
acci0.bs=predict(lm(Fatal~bs(index,df=5+3,intercept=FALSE),data=accident),
                 newdata=accident,interval = 'confidence')
par(mfrow=c(1,2));
plot(accident$index,accident$Fatal,ylim=c(0,90),main='Natural splines');
matlines(accident$index,acci0.ns,lty=c(1,3,3),col=c(4,4,4));
plot(accident$index,accident$Fatal,ylim=c(0,90),main='B splines');
matlines(accident$index,acci0.bs,lty=c(1,3,3),col=c(2,2,2));
```



We can conclude that B splines achieves a better result than natural splines, it has smaller LOOCVSE and MSE.

(b)

The goal is to model the hijacking incidents with fatal accidents using both natural splines and B-splines to fit. We fit the temperature data using natural splines and B-splines with the degree of 3. Various number of inner knots are considered, 0-10.

```
# dataset has been sorted based on the fatal accidents
accident<-accident[sort(accident$Fatal,index=TRUE)$ix,]
```

LOOCV for natural splines and B splines

```
accident.cv.ns1=accident.cv.bs1=rep(0,length(n.knots))
for(j in 1:length(n.knots)){ for(l in 1:length(accident$Fatal)){
  #predict hijacking incidents
  a.ns.pre1=predict(glm(Hijacking~ns(Fatal,df=n.knots[j]+1,
                                intercept=FALSE,Boundary.knots=c(10,80)),
                    data=accident[-l,], family = poisson),
                    newdata=accident[l,],type = "response")
  a.bs.pre1=predict(glm(Hijacking~bs(Fatal,df=n.knots[j]+3,
                                intercept=FALSE,Boundary.knots=c(10,80)),
                    data=accident[-l,],family = poisson),
                    newdata=accident[l,],type = "response")

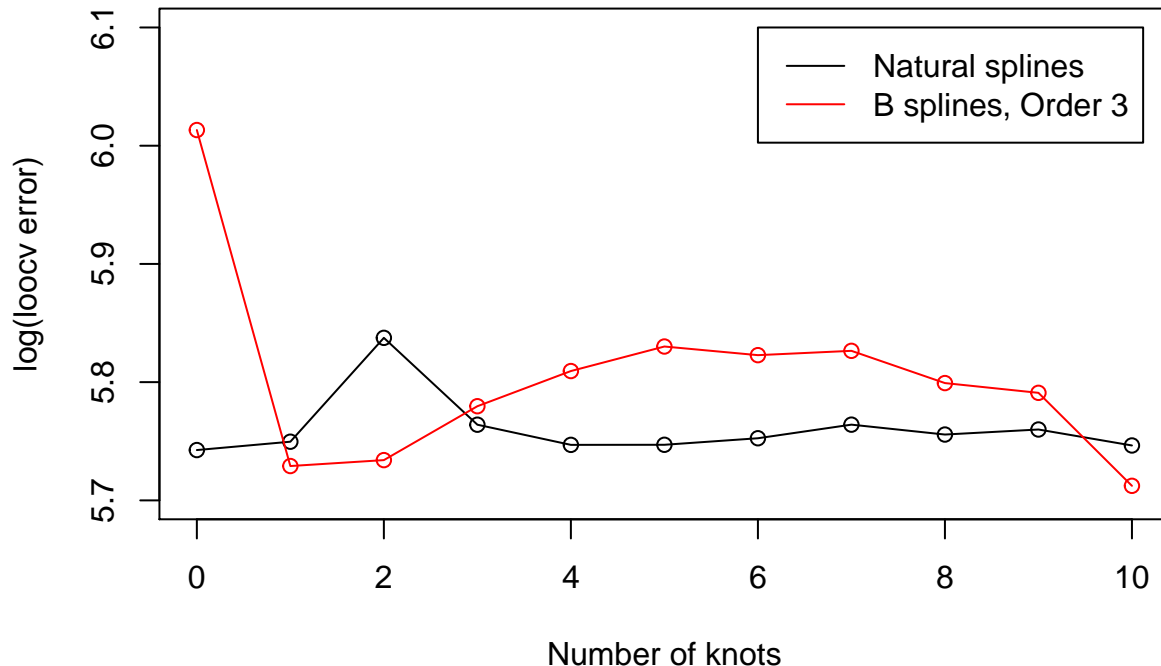
  #cumulative sum of error
  accident.cv.ns1[j]=accident.cv.ns1[j]+(accident$Hijacking[l]-a.ns.pre1)^2
  accident.cv.bs1[j]=accident.cv.bs1[j]+(accident$Hijacking[l]-a.bs.pre1)^2
}}
accident.cv.ns1=accident.cv.ns1/length(accident$Fatal)
```

```

accident.cv.bs1=accident.cv.bs1/length(accident$Fatal)

plot(n.knots,log(accident.cv.ns1),ylab='log(loocv error)',
     xlab='Number of knots','o',ylim = c(5.7, 6.1))
lines(n.knots,log(accident.cv.bs1),col='red','o')
legend(6,6.1,legend=c("Natural splines","B splines, Order 3"),
      col=c(1,2), lty=1)

```



It is not hard to figure out that when the natural splines has the smallest error when number of knots equals 4 and B splines has the smallest error when number of knots equals to 10.

Rather than fitting only natural splines and B splines, we also fit a poisson regression based on that.

```

n.knots.ns <- 4
n.knots.bs <- 10
# Natural splines
acci0.ns1=predict(glm(Hijacking~ns(Fatal,df = n.knots.ns+1,intercept=FALSE),
                     data=accident,family = poisson),
                  type = "response", newdata=accident,
                  interval = 'confidence')

# B splines
acci0.bs1=predict(glm(Hijacking~bs(Fatal,df = n.knots.bs+3,intercept=FALSE),
                     data=accident, family = poisson),
                  type = "response", newdata=accident,
                  interval = 'confidence')

# Poisson model
acci0.p=predict(glm(Hijacking~Fatal,data=accident,family = poisson),
                type = "response",interval = 'confidence')

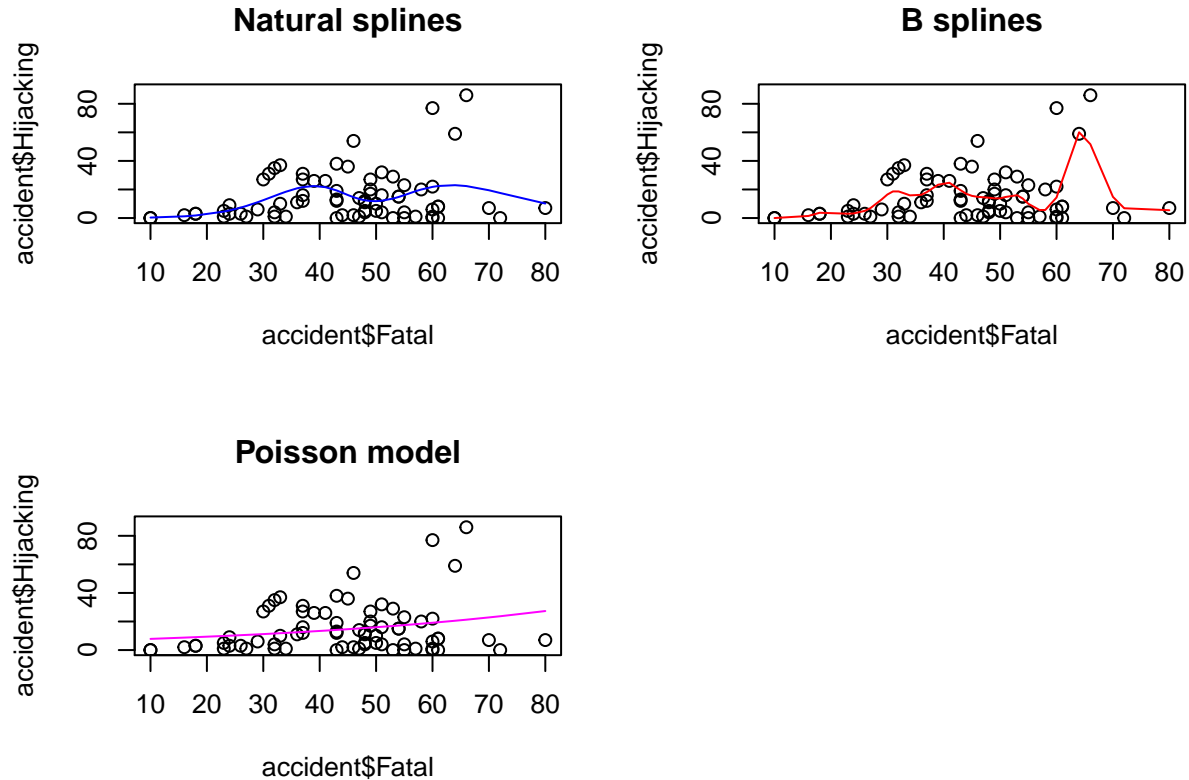
par(mfrow=c(2,2))
plot(accident$Fatal,acci0.ns1,ylim = c(0,90),
     main = "Natural splines")
matlines(accident$Fatal,acci0.ns1,lty=c(1,3,3),col=c(4,4,4))
plot(accident$Fatal,acci0.bs1,ylim = c(0,90),
     main = "B splines")
matlines(accident$Fatal,acci0.bs1,lty=c(1,3,3),col=c(4,4,4))
plot(accident$Fatal,acci0.p,ylim = c(0,90),
     main = "Poisson model")
matlines(accident$Fatal,acci0.p,lty=c(1,3,3),col=c(4,4,4))

```

```

main = "B splines")
matlines(acci0$Fatal,acci0.bs1,lty=c(1,3,3),col=c(2,2,2))
plot(acci0$Fatal,acci0$Hijacking, ylim = c(0,90),
     main = "Poisson model")
matlines(acci0$Fatal,acci0.p,lty=c(1,3,3),col=c(6,6,6))

```



We can conclude that B splines achieves a better result than natural splines, it has smaller LOOCVSE and MSE, the same as question 2(b).

Description of the relationship of those incidents: it is obvious that most of the large Hijacking points concentrates between 25 and 70, Hijacking points are relatively very small in $[10,25]$ and $[70,80]$. There are some very high Hijacking points locating between 60 and 70. Also, most of the Hijacking points locate between $[30,60]$ fatal interval. There can be a potential relationship between Hijacking and Fatal as in $[30,60]$ of Fatal, Hijacking keeps increasing.